

Supplementary Material

LoHR-Bench: A Dual-Level Benchmark for Extended Long-Horizon Robot Manipulation

S1. BENCHMARK TASKS

Tables S1 and S2 summarize the 20 benchmark tasks, grouped into long-horizon manipulation, active exploration, tool use, and cluttered manipulation. Each entry includes the natural-language instruction, number of symbolic predicates (# Sym.), and average episode length in steps. All tasks are evaluated using predicate-based success criteria. Figure S1 shows the corresponding visual settings.

S2. TASK SPECIFICATION

Each task is specified by three files: `problem.pddl`, `domain.pddl`, and `position.json`. Together, these files define the symbolic planning problem, the available actions, and the spatial configuration of the simulated scene.

A. Problem Definition

The problem file declares the objects, the initial state, and the goal as a logical formula over predicates. Goals are expressed as conjunctions of atoms where each atom will be evaluated against the simulated scene to determine task success. Listing S1 presents the `problem.pddl` file for the 10-cube stacking task, in which all cubes are initialized at distinct positions and the goal is to sequentially stack cube B on cube A, cube C on cube B and so forth.

B. Domain Definition

The domain file declares types (e.g., robot, block, surface), predicates (e.g., robot-at, object-at, holding, on-object, clear), and actions (e.g., move-to, pick, place). For the 10-cube stacking example (Listing S2), there are three actions needed to complete the task, i.e., pick, place, and stack. Each action specifies its parameters, preconditions, and effects.

C. Spatial Configuration

The position file specifies spatial regions and initial object placements, and maps position symbols to workspace poses. Task configurations may also specify object properties such as model identifiers, scales, and rotations. For example, in Listing S3, the positions of the cubes and robot are initialized at `pos1-pos11`. The table region is defined by rectangular bounds and a height of 0.02 meters.

S3. PREDICATE EVALUATION

Predicates are defined in `domain.pddl` and evaluated against the simulated state. An object registry maps PDDL object names to scene entities, including robots, objects, and positions. Figure S2 illustrates the verification process, and Table S3 summarizes the predicate checks.

S4. ACTION PRIMITIVES

The experiments use the following action primitives:

- 1) *pick*(*?r*, *?o*, *?p*): it allows the robot *?r* to grasp an object *?o* at position *?p*. The preconditions require that the object is at position *?p* and its top is clear and the gripper is empty.
- 2) *place*(*?r*, *?o*, *?p*): It puts a held object at position *?p*. The primitive only executes if the robot is currently holding an object.
- 3) *place-in*(*?r*, *?o*, *?cont*): It places a held object into a container *?cont*.
- 4) *place-on*(*?r*, *?o*, *?surf*): It places a held object onto another object *?surf*.
- 5) *stack*(*?r*, *?top*, *?bottom*): It stacks the held block *?top* on the clear block *?bottom*.
- 6) *unstack*(*?r*, *?top*, *?bottom*, *?p*): It removes the top block from a stack while the robot is at position *?p*.
- 7) *move-to*(*?r*, *?from*, *?to*): It moves the robot from one location *?from* to another *?to*.
- 8) *open*(*?r*, *?o*, *?drawer*): It opens a drawer by pulling along its axis. Here, *?o* is the cabinet/drawer object in the scene while *?drawer* identifies the drawer.
- 9) *close*(*?r*, *?o*, *?drawer*): It closes a drawer by pushing it in along its axis.

S5. EVALUATION METRICS

We report task success rate, progress score, and completion time.

A. Task Success Rate

Task success rate is the fraction of evaluation trials that achieve a progress score of 1, indicating that all goal atoms hold.

B. Progress Score

For a task with M goal atoms $g_1, \dots, g_M$, the progress score at simulated state s is

$$P(s) = \frac{1}{M} \sum_{i=1}^M \mathbf{1}[g_i(s)], \quad (\text{S1})$$

where $\mathbf{1}[g_i(s)]$ equals 1 when goal atom g_i holds and 0 otherwise. The score lies in $[0, 1]$, and a score of 1 indicates task completion. This metric provides partial credit and can support curriculum learning or reward shaping.

For the 10-cube stacking task, the goal contains nine stacking relations. When three cubes form the correct partial stack, two goal atoms hold, giving $P(s) = (3-1)/(10-1) = 2/9$, as shown in Figure S3.

C. Completion Time

Completion time is the total end-to-end duration of an episode, including planning and action execution, measured using Python’s `time` library.




Fig. S1: Overview of LoHRbench tasks corresponding to Tables S1 and S2.

TABLE S1: Benchmark tasks: long-horizon manipulation and active exploration.

Task	Language instruction	# Sym.	Avg. steps
Long-horizon			
stack_cube_10	Stack 10 cubes together, starting with the red cube.	20	2207
stack_by_color	Group cubes by color and stack each color into its own stack.	20	2213
sort_and_store	Sort tabletop objects into two boxes: toys and balls in one box, tools in the other.	20	2454
reverse_stack	Reverse a stack of 10 cubes.	20	2207
stack_cube_pyramid_simple	Build a 10-cube pyramid with layers 4–3–2–1.	20	2172
stack_cube_hanoi	Solve a Tower-of-Hanoi-style rearrangement with colored cubes using pick-and-place under Hanoi constraints.	20	2127
Active exploration			
stack_cube_active_exploration	Pick up cubes from a closed cabinet and stack them together, starting with the red cube.	10	2092
pick_cube_active_exploration	Pick up the can, screwdriver, and cup from a two-layer drawer and place them at the target location.	10	2134
drawer_search_hard	Pick up the can, screwdriver, and cup from a cluttered two-layer drawer and place them at the target location.	10	2127
stack_cube_cabinet	Pick out several cubes from a closed double-layer cabinet and stack them.	10	1378

TABLE S2: Benchmark tasks: tool use and cluttered manipulation.

Task	Language instruction	# Sym.	Avg. steps
Tool use			
fruit_placement	Place four strawberries into the target position using a container.	10	1349
repackage	Put all cubes into the bowl and stack the bowl on the plate.	7	634
container_transfer	Pour items from one container into another, then place the filled container on a plate.	5	1731
food_storage_simple	Place a strawberry into a bowl, then store the bowl inside a locker.	11	1939
push_T_fruit	Use a T-shaped tool to push multiple strawberries to a target location.	2	240
Cluttered manipulation			
stack_cube_clutter	Stack three cubes together in a compact environment with surrounding obstacles.	18	1516
cluttered_packing	Pick three cubes from a cluttered tabletop and place them into a bowl.	18	1615
tall_clutter	Retrieve a cube surrounded by taller cubes and place it into a bowl.	14	1175
clutter_draw	Extract the red cuboid located beneath the middle column in a three-column stack.	8	828
clutter_push_cube	Push a cube to a target location through clutter, moving obstacles that block the path.	9	983




Fig. S2: Predicate verification using the simulated scene and object registry.




Fig. S3: Partial completion of the 10-cube stacking task: two of nine goal atoms hold, giving a progress score of 2/9.

TABLE S3: Logical predicates and their evaluation criteria.

Predicate	Evaluation criterion
robot-at	The robot’s end effector is within a small distance of the specified position.
object-at	The object is not held and its pose is within a small distance of the specified position.
holding	The robot is grasping the specified registered object.
gripper-empty	The robot is not grasping any registered object.
clear	The top of the specified object or position is clear.
on-object	The vertical separation of the top and bottom objects’ centers and their horizontal offset lie within the defined ranges.
in-object	The object’s position lies inside the specified container region.
in-drawer	The object lies inside the specified drawer volume of the cabinet, with the drawer open.
accessible	A collision-free end-effector approach to a valid grasp pose exists for the object.

TABLE S4: Baseline training configurations. Checkpoint denotes the training iteration used for evaluation.

Method	Batch	Fine-tuning	Checkpoint
ACT	128	Full	50,000
DP	256	Full	30,000
OpenVLA	8	LoRA (rank 32)	200,000
RDT	12	LoRA (rank 32)	150,000
π_0 (FAST)	32	LoRA (rank 32)	100,000

S6. LEARNING-BASED BASELINE DETAILS

A. Shared Observation and Action Setup

We evaluate five baseline policies: Action Chunking with Transformers (ACT), Diffusion Policy (DP), OpenVLA, RDT, and π_0 (FAST). All baselines use the same dataset splits and are conditioned on natural-language task instructions. Observations include RGB images from a base camera and a wrist-mounted camera, together with robot proprioceptive state. Actions are 8-dimensional, comprising seven joint commands and one gripper command. During evaluation, all policies use an execution chunk size of eight actions. This execution chunk size is reported separately from each model’s prediction horizon in Table S5.

B. Training Configuration

All methods use AdamW. ACT, DP, OpenVLA, and RDT use a learning rate of 1×10^{-4} . For π_0 (FAST), the learning rate follows a cosine schedule with a peak of 1×10^{-4} . Table S4 summarizes batch sizes, fine-tuning strategies, and evaluation checkpoints.

ACT and DP are trained end-to-end without parameter freezing and use CLIP (ViT-B/32) for language encoding. OpenVLA, RDT, and π_0 are fine-tuned using Low-Rank Adaptation (LoRA), with backbone parameters frozen and language conditioning handled by their built-in language models. In our experiments, OpenVLA and π_0 use GPUs with 80 GB of memory, while ACT and DP use a 40 GB

TABLE S5: Baseline architectures and prediction horizons. Observation horizons count temporal steps, with two camera views per step; action horizons count predicted actions.

Method	Model	Lang.	Obs.	Act.
ACT	Trans. (CVAE)	CLIP	1	30
DP	Diff. (U-Net)	CLIP	2	16
OpenVLA	VLA (7B)	Built-in	1	16
RDT	VLA (1B)	Built-in	2	64
π_0 (FAST)	VLA (3B)	Built-in	1	16

A100. For RDT, language embeddings are pre-encoded to enable training on a 40 GB A100.

C. Action Chunking with Transformers (ACT)

The Action Chunking with Transformers (ACT) baseline uses a DETR-VAE backbone with a ResNet-18 visual encoder. ACT operates by encoding fixed-length action chunks into a latent representation via a conditional variational autoencoder (CVAE) and decoding them autoregressively conditioned on current visual observations, proprioceptive state, and language instruction embeddings. Language information is provided via CLIP (ViT-B/32) text embeddings that are concatenated with visual and state features prior to decoding. By leveraging chunked action sequences (30 steps), ACT implicitly captures long-horizon behavior while operating on single-step observations. ACT is trained using supervised sequence prediction and optimized with AdamW. The evaluation checkpoint is listed in Table S4. An open-source implementation and reproduction scripts for ACT on LoHRbench are available at ACT and DP reproduction code.

D. Diffusion Policy (DP)

The Diffusion Policy (DP) baseline models action sequences as samples from a conditional denoising diffusion process. A lightweight conditional 1D U-Net processes action trajectories, conditioned on visual observations encoded by a simple CNN, robot proprioceptive state, and language embeddings from CLIP (ViT-B/32). The diffusion process uses a fixed number of denoising steps during training, enabling uncertainty-aware trajectory prediction. DP uses an observation horizon of two timesteps to provide context for multi-step prediction and a prediction horizon of 16 actions. Training relies on DDPM (denoising diffusion probabilistic model) optimization with AdamW, with the evaluation checkpoint listed in Table S4. The reproduction codebase is shared with ACT at ACT and DP reproduction code.

E. OpenVLA

OpenVLA is a large-scale vision-language-action (VLA) model with 7 billion parameters, pretrained on diverse robotic and multimodal datasets. For LoHRbench, we fine-tune OpenVLA using Low-Rank Adaptation (LoRA) to efficiently adapt the model to benchmark-specific tasks while freezing most backbone parameters. Observations consisting of base and wrist RGB images, proprioceptive state, and natural language instructions are fed to the model, which

predicts continuous actions via an L1 regression head. Fine-tuning requires an NVIDIA A100 with 80 GB VRAM due to the model scale. The evaluation checkpoint is listed in Table S4. The OpenVLA reproduction repository is available at OpenVLA reproduction code.

F. RDT

RDT is a diffusion-based vision-language-action foundation model that uses a Diffusion Transformer (DiT) backbone to model multi-modal action distributions, originally developed for complex (including bimanual) manipulation. Given multi-view RGB observations (e.g., base and wrist cameras), robot proprioceptive state, and a natural-language instruction, RDT predicts a fixed-length (64-step) action chunk by denoising in action space. In addition, RDT adopts a unified, physically interpretable action space to better accommodate heterogeneous robot embodiments and action parameterizations. In our LoHRbench setup, we fine-tune RDT under the same observation/action specification as other baselines (8D joint-space control: 7 joint commands + 1 gripper command). In our experiments, we train RDT on a 40 GB A100 and pre-encode the language to reduce on-board GPU memory usage during training and evaluation. We use a fixed learning rate of 1×10^{-4} , and report results using the 150,000-iteration checkpoint for evaluation.

G. π_0 (FAST)

π_0 (FAST) is a vision-language-action model built on a PaLI-Gemma backbone with a FAST tokenizer for discretized action generation. The model receives the same multi-view visual and proprioceptive inputs as OpenVLA, with language instructions encoded internally by the PaLI module. Actions are tokenized into fixed-length (16-step) sequences for autoregressive generation. We fine-tune π_0 on LoHRbench with LoRA, freezing backbone weights and adapting only low-rank layers. Training uses a cosine learning rate schedule with peak learning rate at 1×10^{-4} and terminates at 100,000 iterations. Evaluation uses the final checkpoint. The OpenPI reproduction is available at OpenPI reproduction code.

S7. TAMP BASELINE DETAILS

A. PDDLStream

a) Overview and observability: We implement the PDDLStream baseline as a minimal task-and-motion planning (TAMP) pipeline with a fixed discrete abstraction and open-loop execution. The symbolic initial state is loaded from pre-authored `problem.pddl` files. These files hard-code facts such as `(object-at obj posN)` and `clear`. The baseline therefore uses no perception. In the Active Exploration tasks, objects inside drawers/cabinets are also specified in the initial state. The baseline assumes full observability and does not model discovery or belief.

b) Discrete state representation: We discretize the workspace into named positions `pos1..posN` defined in `position.json`. Objects are assigned to these symbols offline. We do not map continuous poses back to symbols at runtime. If an object ends up at a pose that does not match any `posN`, the symbolic state has no representation for it.

c) Symbolic planning and streams: During planning, PDDLStream calls Fast Downward once to obtain a complete symbolic plan. We use the incremental solver with unit costs (`unit_costs=True`) and a time limit of 300s. Streams only enumerate discrete position symbols. Each generator yields a finite list of `posN` tokens, either hardcoded or enumerated from `(is-position posN)` facts. Stream certification is symbolic. Yielding `posN` certifies facts like `(object-at ?o ?p)` without motion-feasibility checks. Streams do not generate grasps, IK solutions, or trajectories.

d) Motion execution and state updates: In execution, the executor computes all continuous parameters. We use a fixed top-down grasp heuristic with a small set of yaw angles. IK and trajectories are solved by ManiSkill MPlib for all runs. The default primitive `move_to_pose_with_screw` does not perform collision checking, while RRTConnect is used for navigation-style `move-to`. After each action, we re-evaluate predicates from simulator state to update the symbolic state. However, we do not replan. If a motion fails, the episode terminates.

e) Failure modes: These design choices simplify integration and keep the search space finite. They also impose assumptions that break into two common failure modes: (i) cluttered scenes that require explicit rearrangement of blocking objects to make grasps and motions feasible, and (ii) exploration settings where execution must gather information and replan after partial progress.

B. VLM-TAMP

a) Subgoal proposal and privileged context: VLM-TAMP combines a vision-language model (VLM) for high-level subgoal proposal with PDDLStream for low-level planning and motion execution. At each iteration, we query `gpt-4o-mini` with an RGB image and a structured text prompt. The prompt provides the natural-language task goal, the available predicates, and the valid object symbols. Crucially, the prompt also includes a scene description that is *privileged* rather than vision-inferred: the object list and initial relations are hardcoded in PDDL names (e.g., “cubeA at pos1, ..., robot at pos6”). The VLM therefore does not perform object recognition or grounding from pixels; the image is auxiliary context.

b) Semantic checking and planning: The VLM is instructed to output 2–5 intermediate subgoals, and then translate them into formal predicate strings. We parse the response to extract predicate expressions (e.g., `holding(panda_wristcam, cubeB)`). Each candidate subgoal is filtered by a semantic checker that enforces a fixed predicate vocabulary, correct arity, and type-consistent arguments based on the PDDL domain/problem symbols. For each semantically valid subgoal, we inject it as the temporary goal in a new

PDDLStream problem and call Fast Downward to obtain a complete plan that achieves the subgoal. As above, we use the incremental solver with unit costs and a 300s time limit. We then execute this plan in the simulator. Executed actions are appended to an `action_history` string, which is included in subsequent prompts.

c) Retries and termination: VLM-TAMP uses three retry layers. At the API level, failed VLM requests are retried up to 6 times with a linear backoff. At the planning level, if all subgoals in a VLM response fail (semantic check, no PDDLStream solution, or execution failure), we re-prompt the VLM and continue for up to 10 subgoal iterations per attempt. At the episode level, we restart from a fresh environment up to 5 attempts. The method terminates successfully when the environment reports goal satisfaction after plan execution.

S8. EXTENDING LOHRBENCH

A. Adding New Tasks

Adding a new task is relatively straightforward and follows a well-defined, template-based process. A task is specified by creating a directory under `benchmark/tabletop/{task_type}/{task_name}/`, which contains PDDL domain and problem files defining types, predicates, actions, the initial state, and success goals, along with a `position.json` file specifying spatial regions and object placements and a reference symbolic plan. If existing action primitives (e.g., pick, place, move-to) are sufficient, no additional code is required beyond these configuration files. Otherwise, a new task executor can be implemented in `src/tampbench/motion_planner/tasks/tabletop/` to translate symbolic actions into motion commands. Finally, the task is registered by adding an entry to the `TASK_TO_EXECUTOR` mapping in `utils.py`. In practice, most development effort lies in authoring the PDDL specifications, while executor implementation is optional for many tasks.

B. Integrating Other TAMP Methods

Integrating alternative TAMP methods requires moderate effort and is supported by the modular architecture of the motion-planning layer. LoHRbench defines an abstract `BaseMotionPlanner` interface that specifies core motion primitives such as pose reaching and gripper control, with existing implementations including a GPU-accelerated CuRobo backend and a CPU-based ManiSkill/MPLib backend. New TAMP methods can be incorporated by implementing this interface and registering the planner within the factory mechanism used by task executors. Because symbolic task planning is expressed in PDDL and largely decoupled from motion planning, most integration work is confined to the motion layer. Additionally, the multi-level policy structure in `tamp_env.py` allows new methods to be integrated at the task-planning level, the motion-control level, or both.

C. Supporting Other Simulators

LoHRbench is currently tightly coupled to ManiSkill 3 and the SAPIEN simulator, but its design allows extension to other simulators with non-trivial engineering effort. Simulator-specific components include the environment converter that instantiates scenes and objects, the predicate evaluation system that queries grasping and spatial relations, and the motion-planning adapters that interface with simulator state and kinematics. Object loading and asset management are also implemented using ManiSkill-specific pipelines and dependencies. To support alternative simulators such as MuJoCo or Isaac Lab, these components would need to be re-implemented against the corresponding simulator APIs. Importantly, the task specification layer, comprising PDDL definitions, spatial configurations, and symbolic planning, is simulator-agnostic and can be reused directly. A natural extension would be to introduce a simulator abstraction layer, enabling a pluggable PDDL-to-simulator bridge rather than duplicating task logic across backends.

D. Supporting Multiple Robot Configurations

LoHRbench fully supports multiple robot configurations through a unified and extensible infrastructure. This includes a central robot registry (`robot_config.py`), training scripts that accept configurable state and action dimensions, evaluation wrappers that automatically infer dimensions from checkpoints, and a motion-planning layer that adapts to robot-specific parameters such as joint names and CuRobo configurations. Adding a new robot (e.g., UR10e or xArm) requires only registering the robot in `robot_config.py` (approximately 15 lines of code), updating the PDDL files to specify the robot name, and collecting demonstration data. Notably, the robot used during data collection is determined by the PDDL problem file rather than by command-line arguments.

S9. TASK SPECIFICATION LISTINGS

The following listings collect the example files referenced in Section S2. The problem listing is an excerpt; the domain and spatial configuration listings retain the supplied example.

A. Problem File

Listing S1: `problem.pddl` for 10-cube stacking (excerpt)

```
(define (problem StackCube-v2)
  (:domain maniskill-stacking-primitives)
  (:objects
    panda_wristcam - robot
    table-workspace - surface
    pos1 pos2 pos3 pos4 pos5 pos6
    pos7 pos8 pos9 pos10 pos11 - position
    cubeA cubeB cubeC cubeD cubeE
    cubeF cubeG cubeH cubeI cubeJ - block
  )
  (:init
    ;; Robot state
    (robot-at panda_wristcam pos6)
    (gripper-empty panda_wristcam)

    ;; Cube initial placements
    (object-at cubeA pos1)
```

```

...
(object-at      cubeJ pos11)

;; All cubes are clear (can be picked)
(clear cubeA)
...
(clear cubeJ)
)
(:goal
 (and
  (on-object cubeB cubeA)
  (on-object cubeC cubeB)
  ...
  (on-object cubeJ cubeI)
 )
)
)
)

```

B. Domain File

Listing S2: domain.pddl for 10-cube stacking

```

(define (domain maniskill-stacking-primitives)
  (:requirements :strips :typing)

  ;; Types
  (:types
    robot      ; the manipulator
    block      ; manipulated block
    surface    ; tables
    position   ; discrete surface locations
    obj        ; external source
  )

  ;; Predicates
  (:predicates
    ;; Robot state
    (robot-at ?r - robot ?p - position)
    (gripper-empty ?r - robot)
    (holding ?r - robot ?o - block)

    ;; Object state
    (object-at ?o - block ?p - position)
    (on-object ?o1 - block ?o2 - block)
    ;; Can be picked when nothing is on top
    (clear ?x - (either block surface))
  )

  ;; Pick any object at current position
  (:action pick
    :parameters (?r - robot ?o - block
                  ?p - position)
    :precondition (and
      ;; (robot-at ?r ?p)
      (object-at ?o ?p)
      (clear ?o)
      (gripper-empty ?r)
      ;; (reachable ?r ?o)
    )
    :effect (and
      (holding ?r ?o)
      (not (object-at ?o ?p))
      (not (gripper-empty ?r))
    )
  )

  ;; Place any held object at current position
  (:action place
    :parameters (?r - robot ?o - block

```

```

      ?p - position)
    :precondition (and
      ;; (robot-at ?r ?p)
      (holding ?r ?o)
      ;; (clear ?p) ; ensure target is free
    )
    :effect (and
      (object-at ?o ?p)
      (clear ?o)
      (gripper-empty ?r)
      (not (holding ?r ?o))
    )
  )

  ;; Stack a held block onto another block
  (:action stack
    :parameters (?r - robot ?top - block
                  ?bottom - block)
    :precondition (and
      (holding ?r ?top)
      (clear ?bottom)
    )
    :effect (and
      (gripper-empty ?r)
      (not (holding ?r ?top))
      (on-object ?top ?bottom)
      (not (clear ?bottom))
      (clear ?top)
    )
  )
)
)
)

```

C. Spatial Configuration File

Listing S3: position.json for 10-cube stacking

```

{
  "positions": {
    "pos1": {"region": "table-on",
      "fixed_coords": [-0.12, -0.2, 0.02]},
    "pos2": {"region": "table-on",
      "fixed_coords": [-0.02, -0.2, 0.02]},
    "pos3": {"region": "table-on",
      "fixed_coords": [0.08, -0.2, 0.02]},
    "pos4": {"region": "table-on",
      "fixed_coords": [-0.12, -0.1, 0.02]},
    "pos5": {"region": "table-on",
      "fixed_coords": [-0.02, -0.1, 0.02]},
    "pos6": {"region": "robot-at",
      "fixed_coords": [0.08, -0.1, 0.02]},
    "pos7": {"region": "table-on",
      "fixed_coords": [0.08, -0.1, 0.02]},
    "pos8": {"region": "table-on",
      "fixed_coords": [-0.12, 0, 0.02]},
    "pos9": {"region": "table-on",
      "fixed_coords": [-0.02, 0, 0.02]},
    "pos10": {"region": "table-on",
      "fixed_coords": [0.08, 0, 0.02]},
    "pos11": {"region": "table-on",
      "fixed_coords": [-0.2, -0.1, 0.02]}
  },
  "regions": {
    "table-on": {
      "bounds": [[-0.15, -0.25], [0.15, 0.25]],
      "height": 0.02
    },
    "robot-at": {
      "bounds": [[-0.615, -0.615], [0, 0]],
      "height": 0.0
    }
  }
}

```